

Fundraising MCPs (Claude, ChatGPT, & Perplexity) 🚧

One MCP (8Raise) connected to three AI assistants (Claude, ChatGPT, and Perplexity) with HeyReach wired natively through 8Raise. One ranked CSV in. 247 personalized LinkedIn DMs out. 8 calls booked.

8

8FUNDRAISING BY NICO & IGNACIO

APR 29, 2026 • PAID



9



2

Share

Three weeks ago in [the 8Raise Playbook](#), we walked Maya (Flowly's CEO) through an investor-discovery half of a fundraiser. She typed a one-sentence brief, 8Raise ranked 247 named seed VCs at 52 European fintech-adjacent firms, ranked, scored, verified emails and LinkedIn URLs and recent posts. The discovery half took 90 seconds.

Then she sat at her laptop on Thursday morning to actually send the messages.

Forty-seven minutes later she had written six DMs. The first one was great. The second one was good. The third one was a copy of the first one with two words changed. By the sixth one she was rephrasing a sentence she had already used in number two. Her hand-built personalized outreach was already drifting into slop.

Most founders think the bottleneck on a fundraiser is finding the partners. Wrong.

The bottleneck is what happens AFTER you find them. Forty-seven manual DMs in four hours produces six personalized messages and forty-one slop messages. For

slop messages on a fund's LinkedIn inbox is how you get filtered out of every s
the partner ever runs again.

There is a better loop. It has four layers. Each layer is automated. The whole tl
runs in about an hour. We have run it on 23 mandates over the last 12 months.
median outcome on a 50-investor active campaign: 23 connection accepts, 14 i
first calls booked inside two weeks, 3 second meetings inside four weeks. That
lift over manual LinkedIn DMs run by the same founder.

**By the end of this post, you will have a four-layer LinkedIn outreach stack th
named investors with verified profiles, generates personalized messages tied
investor's most recent public post, sends connection requests and follow-ups
scale, and tracks responses back into your investor pipeline. Three AI assist
supported (Claude, ChatGPT, and now Perplexity since their custom-connec
launched). One paste-ready Claude project plus the system prompt, the Link
rate-limit policy, and the response-handling loop. About one hour of setup.**

Here is what is inside.




Table of contents

1. Why manual LinkedIn outreach for fundraising is broken
2. The four-layer stack
3. Layer 1: investor discovery (the input)
4. Layer 2: the AI assistant that writes the messages
5. Layer 3: the outreach automation that sends them
6. Layer 4: the response handling loop
7. Worked example: Maya's Q2 European fintech campaign
8. Connection setup (the three integrations)
9. Compliance, rate limits, and what NOT to do
10. Bonus: paste-ready CLAUDE.md plus the message templates

👁️ Before we get into it, a quick heads up: we raise the subscription price ev new subscribers.

If you have been thinking about upgrading to premium, now is the moment. C subscribers lock in today's rate forever.

After the next 100 subscribers, new subscribers will pay significantly more. No just a real counter.

The reason? We are working on 100xing the value of this publication. You are love what is coming next 🔥

● Prices increase after the next 100 subscribers, subscribe today to lock in your rate permanently

8

8fundraising

Subscribe now.
Pay less, forever.

After the next 100 subscribers, prices go up significantly. Lock in today's rate for life.

64
SUBS

:

7
DAYS

:

0
HRS

:

0
MIN

PLAN		AFTER NEXT BUYER
Monthly	EUR 20/mo today	EUR 29/mo +45%
Annual	EUR 150/yr today	EUR 249/yr +66%

Subscribe at today's price →

No contracts · Cancel anytime · Rate locked for life

€20/month or €150/year. Founding Member ("F&M + AI Tools") at €240/year for access to our tooling plus a live pitch review.

<https://8fundraising.substack.com/p/fundraising-mcps-claude-and-chatgpt>

Page 4 of 39

Inside you get:

- 10+ AI templates across fundraising, marketing, and operations
- 150+ pitch decks that have raised capital
- 3,000+ investors and accelerators database
- Every future playbook in this series

Remember!

As a premium subscriber you also get access to every other playbook in this series such as:

- 1 [The AI Investor Scraping Playbook for Claude](#)
- 2 [The X Posting Playbook for Claude](#)
- 3 [The 8fundraising Investor List](#)
- 4 [The Best Claude Connectors for Raising Capital](#)
- 5 [The Pitch Decks Playbook](#)
- 6 [The Pitch Decks Playbook for Claude](#)
- 7 [The LinkedIn Playbook for Claude: Profile + Posting](#)
- 8 [The Investor Outreach Autopilot for Claude](#)
- 9 [How to Scrape Investors with AI Properly](#)

10 [How to Create Landing Pages Like a Pro](#) (just shipped)

1 1 (you are here)

🗨️ Or you can always invite a friend and unlock rewards along the way:

→ 3 friends = 1 month free

→ 5 friends = 3 months free

→ 25 friends = 6 months free

Early founders and operators are already inside.

🚀 Let's begin!

Quick note on how to use this post. Copy and paste this entire article into a Claude conversation and ask Claude to walk you through the loop on your own computer. This post is written to be runnable, not just readable. Claude is the guide. You are the founder answering its questions.

🔧 **The two tools you'll wire into this loop:**

→ [8Raise](#) — the investor discovery layer (Layer 1). Plain English brief in, ranked partners with verified emails and LinkedIn URLs out, in 90 seconds.

→ [HeyReach](#) — the LinkedIn outreach automation layer (Layer 3). Sends connection requests and follow-up DMs from your real LinkedIn account at LinkedIn-safe limits, with response tracking back into a single dashboard.

Open both in a tab now. Setup happens in Section 8.

1. Why manual LinkedIn outreach for fundraising is broken

Manual LinkedIn DMs at fundraising scale fail for three structural reasons. None of them is “you are bad at writing.”

Reason 1: personalization decays. The first three DMs you write on a Thursday morning are sharp. The fourth feels familiar. By DM number eight you are reusing the same opening sentence with two words swapped. By DM number twenty you are writing the same DM. The decay is not laziness. It is cognitive load. The human runs out of unique opening hooks somewhere between message 5 and 8 even when you have detailed context on the recipient.

Reason 2: timing is asymmetric. Investor X posts something interesting on Thursday at 11 AM. The “reference your recent post” hook is golden for ~36 hours. By Friday afternoon when you finally write the DM, the post is buried, the reference feels stale, and the partner reads “I saw your post about [X]” and quietly assumes you do not actually read LinkedIn (because real readers reference posts within hours, not days).

Reason 3: response handling kills the funnel. Six partners reply. You spend 90 minutes responding to all six because you have been doing other work in between. Two of them go cold because the reply window passed. Two of them go cold because your response was rushed. Two of them book calls. The conversion of “interested reply” to “booked call” should be 60 to 70%. Manual fundraising founders run at ~35%. The gap is response velocity.

The fix is not “write better DMs.” It is to wrap the four-layer loop below around your work so the human is the judgment layer, not the typing layer.

Rule 1: discovery feeds personalization. The output of layer 1 is a CSV with n

partners, LinkedIn URLs, and (critically) each partner's last 3 to 5 public posts. The personalization layer needs the posts as input. If your investor list does not have posts attached, you are setting the AI up to write generic.

Rule 2: every message has a single hook. One post per partner. One reference message. Multi-hook DMs read like LinkedIn-marketing-Twitter and convert less than single-hook DMs across every dataset we have measured.

Rule 3: human approval gate before send. No DM auto-sends without your real-time approval. Message Generation is automated. Sending is approved. This rule is what keeps you from accidentally messaging the same partner twice or DMing someone who left the list last week.

Everything below follows those rules.

(INSERT IMAGE HERE: three-rule diagram showing personalization decay, time-to-response asymmetry, and response velocity as the three failure modes, with the four-layer stack on the right | post-08-failure-modes.png)

2. The four-layer stack

Four layers. Each one is automated. The whole loop runs in about one hour of setup time and then runs in the background for 14 days.

Layer 1: discovery. Input: a one-sentence brief about who you are raising from. Output: a ranked CSV of named investors with verified LinkedIn URLs, verified emails, and each partner's last 3 to 5 public posts attached. We use 8Raise for this (covered in [Post 6](#)).


Layer 2: personalization. Input: the CSV from Layer 1. Output: 50 to 250 personalized

opening DMs, each tied to a specific recent post by that partner. This is the AI assistant layer. Claude, ChatGPT, and Perplexity all work here. We use Claude internally because of the way it follows brand-voice rules across 250 outputs in ChatGPT works equally well for message generation. Perplexity adds live web grounding which sharpens hooks when the partner posted breaking news inside last 48 hours. The post below covers all three.

Layer 3: delivery: Input: the 50 to 250 personalized DMs from Layer 2. Output: connection requests sent on LinkedIn, follow-up DMs queued, daily rate limit respected, all tracked back into a single dashboard. We use HeyReach for this with a native 8Raise → HeyReach push that sends search results straight into a running campaign in one click. No CSV juggling, no copy-paste. Setup is one-time, takes ~60 seconds.

Layer 4: response handling. Input: the replies that come back. Output: classified responses (interested, not now, wrong fit, asked a question), tagged in your pipeline with a draft reply queued for each one. The same AI assistant from Layer 2 handles this. The human approves and sends.

Run end to end, the median time from “I want to send 50 DMs” to “the campaign running” is about 60 minutes. The DMs themselves go out over 14 days because LinkedIn rate limits do not let you send 50 at once. We will get to the rate limit section 9.



3. Layer 1: investor discovery (the input)

The discovery layer is 8Raise. We covered the full 8Raise playbook in [Post 6](#), so this section focuses specifically on how to prompt 8Raise for the LinkedIn outreach. The five rules that determine whether your search returns 50 dialed-in partners or mostly-dead leads.

Rule 1: Describe the investor profile, not the company you're raising for. 8Raise is to match a target ICP. The more concrete the ICP, the tighter the result set.

→ Weak: "I'm a fintech startup raising a seed round, find me investors"

→ Strong: "Find me seed-stage VCs in the UK, Germany, and the Netherlands \$500K to \$2M checks into B2B fintech, especially payments and embedded fin

want partners and principals.”

The strong version gives 8Raise four signals (stage, geography, check size, sector). The weak version forces clarifying questions before anything runs.

Rule 2: Geography. Be specific or be regional, never both vaguely. 8Raise respects geography city → state → country → continent. Mixing levels works if you know what you’re doing (“VCs in Western Europe focused on Germany, France, the Netherlands”). Mixing them by accident kills recall (“VCs in Europe AND London” — Europe includes London, so the query collapses).

Rule 3: One main sector. Others as flavor. 8Raise phrase-matches against firm descriptions for the primary sector signal. “Fintech VCs with a side interest in healthtech” returns a tight result. “VCs that do fintech, healthtech, climate, and proptech” returns near-zero matches because no firm self-describes as all five.

Rule 4: Sample before you go big. Run a 10 to 15-lead test query first to verify filters look right. Then crank to 50 for the LinkedIn campaign. Asking for 100 on the first query is the single most expensive habit we see across mandates.

Rule 5: Use the confirmation block, every time. Before any search runs, 8Raise gives you a filter preview (mode, sectors, geography, exclusions, lead count). This is how you catch mistakes for free. Adjust before you say “run”:

→ “actually include hedge funds” → drops Hedge Fund from the exclude list

→ “narrower precision, US only” → tightens search precision and geography

→ “skip phone, just email and LinkedIn” → kills phone enrichment (faster results)

Saying “run” before reviewing the preview is the single biggest preventable cost of the product.

The LinkedIn-outreach prompt template (copy, fill, run):

Find me 50 [STAGE] VCs in [GEOGRAPHY] investing in [SECTOR], who write checks of [CHECK SIZE] and have done deals in the last 12 months. I want partners and principals.

For each row, also pull the partner's last 3 to 5 public LinkedIn posts (full text), the timestamp of the most recent one, and a recommended single-post hook based on the most recent post. Export with the post text and hook columns included in the CSV.

Email + LinkedIn, skip phone. Deliver as CSV.

Filled for Maya's Q2 European fintech raise:

Find me 50 seed-stage VCs in the UK, Germany, France, and the Netherlands investing in B2B fintech, especially payments and embedded finance, who write checks of \$500K to \$2M and have done deals in the last 12 months. I want partners and principals.

For each row, also pull the partner's last 3 to 5 public LinkedIn posts (full text), the timestamp of the most recent one, and a recommended single-post hook based on the most recent post. Export with the post text and hook columns included in the CSV.

Email + LinkedIn, skip phone. Deliver as CSV.

That single addition (the “post text + hook” columns) turns the standard ranke output into a personalization-ready dataset. The CSV that comes back has the columns per row:

- partner name, fund, role, score, LinkedIn URL, verified email
- post 1 text, post 1 timestamp, post 1 hook recommendation
- post 2 text, post 2 timestamp, post 2 hook recommendation
- post 3 text, post 3 timestamp, post 3 hook recommendation
- partner-fit score axes (sector, geo, stage, recent activity)

That's the input for Layer 2.

Cost: 50 leads off your 8Raise plan for one 50-investor LinkedIn campaign. Plays from Post 6 apply. Time: about 90 seconds for the search, another 30 seconds to enrich with the hook columns enriched.

If you're not running VC outreach (e.g. you're a fund manager raising from LPs, a RE operator raising for a deal), the LP-mode and Real-Estate-mode prompts, plus the refinement prompts and the saved-search re-run workflow live in [Post 8Raise Playbook](#).



4. Layer 2: the AI assistant that writes the messages

The personalization layer is where Claude or ChatGPT turn the CSV into 50 personalized opening DMs.

Both work. Pick whichever you already use day-to-day. The decision matrix:

Use Claude if: you already have Claude wired into your workflow (which you have been following this series). Claude follows the brand-voice rules and outputs in a row better than ChatGPT in our testing. The bigger context window also means you can drop the entire 50-row CSV into one conversation and it generates all 50 messages in a single pass.

Use ChatGPT if: you already pay for ChatGPT Plus or Pro and you do not want to add a second AI subscription. ChatGPT's GPT-5 (or whatever the current top model is when you read this) is faster on per-message latency. The trade-off is that you have to chunk the CSV into batches of 10 to 15 to avoid output drift.

Use Perplexity if: you want live web-search grounding baked into the personalization layer. Perplexity Pro, Max, or Enterprise (free tier does not have the custom-connector option) layers fresh public-source context onto each message before generating the DM, which makes a real difference on time-sensitive topics like a partner who posted news in the last 48 hours, a fund that just announced a deal, a public hire that closed yesterday). Slower than Claude on bulk generation, but the sharpest message quality on news-anchored hooks.

Mix any two or all three if: you want to A/B test message style. Generate the top 50 DMs across two or three hosts, hand-pick the best 50. We have run this comparison on five mandates and the picked-best-of-multiple campaign consistently comes out about 12% better than any single host alone. This is the one we use. It works in both Claude and ChatGPT (paste it as the first message in Claude, paste it into the "system instructions" field in ChatGPT).

You are the LinkedIn DM writer for an 8fundraising client raising capital from venture investors. You will receive a CSV row containing a partner's name, fund, role, recent LinkedIn posts, and a recommended hook. Your job: write a single 280-character LinkedIn DM that:

1. Opens with a SPECIFIC reference to the recommended hook (the most recent post). Reference one detail from the post that proves you read it. Do NOT say "I saw your post about X" generically.
2. Pivots in one sentence to what the founder you represent does, with one concrete number (€XXX MRR, growth multiple, named

- customer, etc.). The pivot must logically follow from the hook.
3. Closes with a low-friction ask. Either a 15-minute call, a deck share, or a single question. Never a hard ask in a first DM.
 4. Total length: under 280 characters. Including the salutation.
 5. Tone: direct, technical, peer-to-peer. Never "I hope this finds you well." Never "I would love to connect." Never "leverage," "unlock," "synergy." Sentences are short. One idea per sentence.
 6. NEVER use em dashes. Use commas, colons, or periods.

Output format: a single message string, nothing else. No commentary.
No "here is your DM." Just the message.

If the recommended hook is older than 7 days, use the second-most-recent post instead. If all 3 posts are older than 7 days, use the partner's fund-level news (the firm's most recent investment, hire, or memo) as the hook instead.


Paste the CSV row, get a DM. Repeat 50 times (or paste 50 rows and have the AI go through them).

For Maya's worked example (section 7 below), the first DM the AI generates as partner Ana García López's most recent post (a thread about Spanish fintech underserving SME segments, posted 4 days ago) reads:

Ana, your thread on SME-segment underserving in Spanish fintech mirrors we saw building Flowly. Mid-market SaaS teams sit in the same gap (47 integrations, no engineering team to wire them). €280k MRR from 18 customers since launch. Would a 15-min walkthrough be useful?

273 characters. Specific reference. Concrete number. Low-friction close. Ready to send.

Iterate twice on the prompt if the first 5 DMs feel generic. Most prompt issues up by message 3. If they do not show by message 5, the prompt is working.



5. Layer 3: the outreach automation that sends them

Layer 3 is HeyReach. HeyReach is a LinkedIn outreach platform built for sales that automates connection requests, follow-up DMs, and response tracking, all respecting LinkedIn's rate limits.

The setup. You connect HeyReach to your LinkedIn account (the one belonging to the founder doing the outreach, never a fake or scraper account). HeyReach uses your LinkedIn session, sends messages from your account, and respects the daily limit.

LinkedIn imposes on real users.

Two ways to get leads into HeyReach: the native 8Raise → HeyReach push (recommended, ~30 seconds per batch) or the manual CSV upload (the old way minutes per batch). Path A is what we run today. Path B still works if you have connected 8Raise to HeyReach yet.

Path A: native 8Raise → HeyReach push

One-time HeyReach setup (under 60 seconds):

→ In HeyReach (paid tier, free does not expose API keys), left sidebar → Integrations → HeyReach API → Generate. Copy the key once. It is only revealed once.

→ In your 8Raise dashboard → Outreach → paste the key → Connect HeyReach. The key is stored server-side, never exposed to your browser. Disconnect from the page or settings to delete it instantly.

Inside any AI host (Claude, ChatGPT, or Perplexity) or the 8Raise dashboard, a search returns: type “send these to HeyReach” or click the Send to HeyReach button. The picker lists your active campaigns and leads-lists. Pick a campaign, leads chunks of up to 100, deduped on HeyReach’s side.

Two destinations:

→ **Campaign (outreach starts immediately)**: pushes into a running ACTIVE campaign. The campaign’s connect-request + follow-up sequence kicks off as soon as the lead lands. Requirement: the campaign must be ACTIVE in HeyReach. PAUSED, FINISHED, and DRAFT campaigns are greyed out in the picker.

→ **Leads list (holding pen, no outreach yet)**: pushes into a HeyReach Leads list with no campaign assigned. Nothing gets messaged until you move them into a campaign.

on HeyReach's side. Useful for staging while you build a sequence, or pushing when you do not have an ACTIVE campaign yet. You can create a new list in the Send modal.

Standard workflow with the native push:

1. In HeyReach, create your fundraising campaign once. Configure the 4-step sequence: → Step 1 (day 0): connection request with a 280-character note (personalized DM from Layer 2). → Step 2 (day 3, only if connection accepted) follow-up DM expanding on the pitch with one concrete number and a calendar link. → Step 3 (day 7, only if no reply to step 2): a single short DM with a data angle (e.g. a customer logo or a partner-relevant data point). → Step 4 (day 10 if still no reply): a single-line "no worries if not a fit, will close the loop" note.
2. Set rate limits. Default: 25 connection requests per day, 50 follow-up DMs randomized hours, weekday-only. These are LinkedIn's safe-zone numbers; exceeding them gets your account flagged.
3. Inside Claude / ChatGPT / Perplexity, run your 8Raise search. After the AI returns the personalized DMs from Layer 2, type "send these to HeyReach." Pick 1 campaign. Leads land. The sequence fires automatically.
4. Watch the dashboard. HeyReach tracks: invites sent, accepts, replies, calendar bookings (if you wire your Calendly into the sequence). Refresh once per day.

Path B: manual CSV upload (the old way)

If you have not yet connected 8Raise to HeyReach, the original workflow still works.

1. Create a campaign in HeyReach.
2. Upload the CSV from Layer 1. Map the LinkedIn URL column.
3. Paste the 50 personalized messages from Layer 2. HeyReach has a CSV import tool.

this exact use case.

4. Configure the 4-step sequence (same as Path A).
5. Set rate limits (same as Path A).
6. Launch.

Path B takes 10 to 15 minutes per campaign. Path A takes ~30 seconds. Once you have the native integration set up (60 seconds, one-time), you never run Path B again.

For Maya's Q2 campaign (run on the native push): 50 connection requests, sent over 5 working days. 23 accepts (46%, above the 30 to 40% benchmark). 14 of the accepts replied to the day-3 follow-up (61%, above the 40 to 50% benchmark). 8 first calls booked. 3 second meetings within four weeks.

Cost: HeyReach starts at €79/month for one LinkedIn account, with the Starter plan covering up to 1,000 sends per month. For a single founder running outreach for their company, this is the right plan.



6. Layer 4: the response handling loop

Replies start coming back on day 1. The response handling layer is what turns “interesting reply” into “booked call” without burning your Friday afternoon.

Two integrations make this work. One is the same AI assistant from Layer 2 (Claude or ChatGPT) classifying and drafting responses. The other is HeyReach’s reply notification flow, which can pipe replies into a Slack channel, a webhook, or a database.

The workflow:

1. **HeyReach detects a reply.** It tags the lead as “replied” and pushes a notification to Slack (or wherever you wired it).

2. **You forward the reply text to Claude / ChatGPT** with the context of the c DM and the partner's profile. Use this prompt:

You are the response handler for an 8fundraising client. The partner [NAME] at [FUND] just replied to our outreach DM. Below is our original DM, their reply, and their profile context.

Classify the reply as one of:

- INTERESTED (wants more, asking for a call or deck)
- INTERESTED-LATER (will engage in a defined window)
- WRONG-FIT (clearly not their thesis)
- WRONG-NOW (timing issue, but might fit later)
- QUESTION (asked a specific question first)

Then draft a single response under 320 characters that:

- Matches their energy level (do not over-pitch a soft yes)
- Includes the next concrete step (calendar link, deck share, one-line answer to their question)
- Does NOT use "I appreciate you taking the time" or any other filler
- Stays in the same brand voice as the original DM


Output: classification on line 1, draft response on line 2. Nothing else.

3. **Read the draft. Edit if needed. Send.** Total time per response: about 90 seconds including the read.
4. **Update the pipeline.** The classification flows into your Notion or Attio pipeline (using the same Composio-driven setup we covered in [Post 4](#)). The lead's status updates from "replied" to its specific classification.

5. For INTERESTED responses, the calendar link is in the draft already. The partner books or does not.

For Maya's Q2 campaign, the 14 replies broke down as: 6 INTERESTED (booked 3 INTERESTED-LATER (added to a 90-day re-engagement queue), 2 WRONG (closed out, tagged "not a fit, q2 fintech"), 2 WRONG-NOW (added to "check back q4"), 1 QUESTION (about Flowly's compliance posture, answered with a one-paragraph reply and a follow-up call booked).

Total response handling time across 14 replies: about 25 minutes over 14 days.



7. Worked example: Maya's Q2 European fintech campaign

Pulling all four layers together. Maya at Flowly is running her Q2 fundraising prep. She wants 50 active investor conversations by end of June.

Tuesday morning. Maya opens Claude with 8Raise wired in. She types: “8Raise me 50 seed VCs in Europe investing in fintech-adjacent B2B SaaS, check size > €2M, currently active in 2026, partner-led firms, contact via email or LinkedIn each row, pull the last 3 LinkedIn posts plus a recommended hook.” 8Raise returns partners, scored. She filters to top 50 by score, exports the CSV with the post text and hook columns.

→ Time elapsed: 4 minutes.

→ Cost: 50 leads off her 8Raise plan.

Tuesday morning, 11 AM. Maya pastes the system prompt for Layer 2 into Claude. She drops the 50-row CSV into the conversation. Claude generates all 50 personalized DMs in one pass, each one referencing a specific post by that partner. Maya reviews through them, edits 4 (where the AI’s hook felt forced), and approves the rest.

→ Time elapsed: 22 minutes.

→ Cost: ~€0.40 in Claude API spend.

Tuesday afternoon, 2 PM. Maya opens HeyReach. She creates a new campaign, uploads the 50 leads, pastes the 50 messages, and configures the 4-step sequence: connection request, day 3 follow-up, day 7 different angle, day 14 close the loop. Limits at the safe LinkedIn defaults.

→ Time elapsed: 18 minutes.

→ Cost: €79/month for HeyReach Starter.

Tuesday afternoon, 2:18 PM. Campaign is live. Total active setup time: 44 min from “I want to run a campaign” to “the campaign is running.” Add another 16 minutes of email and Slack reading and we are at the round one-hour mark.

The next 14 days. HeyReach rolls out the sequence. Maya checks the dashboard per morning (about 90 seconds). Replies come in, she runs them through the re-handler prompt, sends drafts after a quick edit. Total response-handling time days: 25 minutes.

Outcomes after 14 days:

- 50 connection requests sent.
- 23 accepted (46%).
- 14 replied to the day-3 follow-up (61% of accepts).
- 8 first calls booked.
- 3 second meetings booked inside the next two weeks.

Compare against Maya’s previous fundraise (manual LinkedIn DMs, no automation). The same 50-investor target list took 18 hours of writing time, generated 27 connection requests (because she ran out of stamina at 27), 9 accepts, 3 replies, 1 first call. The four-layer stack delivered an 8x lift on calls booked at roughly 7% of the time investment.

8. Connection setup (three platforms, same OAuth flow)

The whole stack works in Claude, ChatGPT, or Perplexity. All three web paths use OAuth, which means zero API keys to copy on the web for any of them. You sign up for 8Raise once in a browser window, click Allow on the consent screen, and you're done. The Desktop and Claude Code paths still use the Bearer token if you prefer to use the terminal.

The values you will paste in either flow:

→ MCP connector name: 8Raise

→ MCP server URL: <https://mcp.8raise.com/mcp>

That is it. Same two values for every path.

Path A: Claude.ai (web), about 60 seconds

1. Sign in to claude.ai. In the left sidebar at the top, click **Customize**.
2. On the Customize page, click **Connectors** in the left sub-nav inside the “Connect your apps” card.
3. In the Connectors panel, click the + button at the top right of the connect section.
4. Choose **Add custom connector** from the popover (the alternative is “Browse connectors”, which is the marketplace; we want custom).
5. Paste **8Raise** into the Name field and <https://mcp.8raise.com/mcp/8raise> into the Remote MCP server URL field. Leave Advanced settings as-is. Click **Add**.
6. Claude opens a sign-in window pointing at 8raise.com. Sign in (or sign up if you don't have an account). Click **Allow** on the consent screen. Done.

The connector now lives in your Claude.ai sidebar. Type **Hi 8Raise** and the flow kicks in.

Path B: ChatGPT (Pro / Plus / Business / EduPro+), about 60 seconds

ChatGPT's custom-MCP feature requires Developer Mode, which is only available on the paid tiers above. Free accounts cannot add custom MCPs as of mid-2026.

1. Sign in to chatgpt.com. In the left sidebar, click **More** → choose **Apps** from the popup.
2. On the Apps page (Settings → Apps), click **Advanced settings** on the right side of the page.
3. Toggle **Developer mode** ON. ChatGPT will warn you it is experimental. Type **Hi 8Raise** and the flow kicks in.

expected. You are loading your own server, which you trust. Leave the “Er CSP” toggle on as is.

4. Click **Back** to return to the Apps page. Click **Create app** on the right of the Advanced settings row.
5. Paste **8Raise** into Name and **https://mcp.8raise.com/mcp** into MCP Server URL. Authentication: leave as **OAuth** (the default). Tick the “I understand and want to continue” checkbox. Click **Create**.
6. ChatGPT opens a sign-in window pointing at 8raise.com. Sign in (or sign up) and click **Allow** on the consent screen. Done.

The connector now appears in ChatGPT’s app sidebar. Type **Hi 8Raise** and the welcome flow kicks in.

Path C: Perplexity (Pro / Max / Enterprise), about 60 seconds

Perplexity’s custom-connector feature requires a Pro, Max, or Enterprise account. The free tier does not expose the Custom connector option. On Enterprise, an admin needs to enable custom connectors first (Enterprise settings → Permissions → Connectors).

1. Sign in to perplexity.ai. In the left sidebar, click **Customize** to expand the customization sub-nav.
2. Click **Connectors**.
3. On the Connectors page, click **+ Custom connector** in the top-right.
4. In the modal, select **Remote** at the top. Paste **8Raise** into Name and **https://mcp.8raise.com/mcp** into MCP server URL. Open Advanced settings and set Authentication = **OAuth** (leave Client ID / Secret blank — our server signs

OAuth discovery, so Perplexity fills these automatically). Tick the acknowledgement box. Click Add.

5. A sign-in window opens. Sign in to 8Raise, click Allow. The connector lig your Connectors list.

The connector now surfaces inside Perplexity Computer workflows and chat. searches use the same background-polling UX you would see on ChatGPT — Perplexity says “running” and checks back as results land.

Path D: Claude Desktop or Claude Code (terminal user Bearer token path)

For founders who run Claude in the desktop app or in the CLI, the Bearer token is faster than the OAuth web flow.

1. Sign up at 8raise.com, open Settings, generate your API key. The key starts with 8r_ and is shown once. Copy it.
2. For Claude Desktop, paste the following block into `~/Library/Application Support/Claude/claude_desktop_config.json` (macOS) or `%APPDATA%\Claude\claude_desktop_config.json` (Windows), replacing `YOUR_8RAISE_KEY`:

```
{
  "mcpServers": {
    "8raise": {
      "url": "https://mcp.8raise.com/mcp",
      "headers": { "Authorization": "Bearer YOUR_8RAISE_KEY" }
    }
  }
}
```

For Claude Code, run:

```
claude mcp add --transport http 8raise https://mcp.8raise.com/mcp \
  --header "Authorization: Bearer YOUR_8RAISE_KEY"
```

3. Restart Claude Desktop (or open a new Claude Code session). The connection with all 16 tools registered.

Connecting HeyReach (the new way: one API key in the 8Raise dashboard)

HeyReach used to require its own MCP in your AI host. As of 04/30/2026, the new 8Raise → HeyReach integration replaces that flow. You set up HeyReach once in the 8Raise dashboard, and every search you run in Claude / ChatGPT / Perplexity gives you a “Send to HeyReach” action automatically.

One-time setup, ~60 seconds:

1. In HeyReach (paid tier required), left sidebar → Integrations → HeyReach Generate. Copy the key.
2. In 8raise.com/dashboard → Outreach → paste the key → Connect HeyReach
3. Done. The “Send to HeyReach” picker now appears in every 8Raise search

Note: the old MCP-based HeyReach setup (adding HeyReach as a separate connection in your AI host) is no longer needed. If you set it up before, the new native push button supersedes it. You can leave the old MCP installed without harm; it just is not the recommended path anymore.

HeyReach also requires its Chrome extension installed (separately from the AI

actually drive your LinkedIn session at delivery time. Install it once, log in to it in the same browser, and HeyReach reads the session.

9. Compliance, rate limits, and what NOT to do

LinkedIn does not love automation, but it does not punish well-behaved authors either. The four-layer stack stays inside the lines. Going outside the lines gets your LinkedIn account suspended, and a suspended account during a fundraiser is a week problem.

The hard rules:

→ **Never exceed 25 connection requests per day.** This is LinkedIn's invisible soft limit for paid accounts (lower for free accounts, around 100 per week). HeyReach

default rate limits respect this.

→ **Never run two LinkedIn automation tools at once.** If you have HeyReach running and you also use a Sales Navigator automation Chrome extension, LinkedIn flags your account for “scripted activity.” Pick one tool. Disable everything else.

→ **Never send the same DM to two partners.** This is what burns founders. Even if the AI is generating, manually verify that the first 10 messages are clearly distinct. If any of them have identical opening sentences, regenerate the prompt.

→ **Never auto-send. Always have a 4-eyes review.** Layer 2 generates. You read and approve. Then HeyReach sends. The 90 seconds you spend reading the AI’s output is what catches the one DM with the wrong fund name or the wrong partner.

→ **Never reuse a campaign list more than every 6 months.** If a partner did not respond in your Q2 campaign, do not re-message them in Q3 with a different message. Wait at least 6 months. LinkedIn fingerprints repeated outreach to the same partner and increasingly down-ranks it.

→ **Never DM partners who have left the firm.** This is the most common error. Your CSV is a snapshot. By week 2 of your campaign, 1 to 3 of the 50 partners may have moved firms. Run an 8Raise refresh on day 7 of the campaign and pull the 50 LinkedIn URLs through the partner-rotation check. Pause sequences for any partner who has rotated.

→ **Never message LP-class allocators on LinkedIn cold.** This is a fundraising rule: family offices and institutional LPs heavily prefer warm intros and email. LinkedIn cold DMs to LPs convert at <2% and burn the relationship. For LP outreach, use the LP-mode workflow from Post 6 with email + warm-intro paths only.

The worst-case scenarios when you break these rules: 7-day LinkedIn restriction

cannot send any DMs for a week), 30-day suspension, or in the worst case per account ban. We have seen all three across founders we have advised. None of recoverable mid-fundraise.

10. Bonus: paste-ready CLAUDE .md plus the message templates

Drop the file below at the root of a fresh Claude Code project. Wire 8Raise and HeyReach MCPs. Type “read CLAUDE.md and confirm the loop.” Claude takes it from there.

```
# LinkedIn outreach operator: Claude project instructions
```

Role

You are the LinkedIn outreach operator for an 8fundraising client. You run the four-layer loop: discovery, personalization, delivery, response handling. You do NOT auto-send DMs. You generate and queue. The founder approves before HeyReach delivers.

Connectors available

- 8Raise MCP (investor discovery, ranked CSV with LinkedIn posts, native HeyReach for delivery)
- (Optional) Composio MCP for Notion / Slack / Gmail integrations

Files in this project

- /facts/company-facts.md: positioning, audience, traction numbers
- /facts/founder-profile.md: who is doing the outreach, their voice
- /searches/: saved 8Raise search briefs, one per campaign
- /campaigns/: campaign-specific files (DMs, replies, status)
- /reviews/: approved DM batches, dated

The four-layer loop

1. DISCOVERY (8Raise)

- Run the 8Raise search with the founder's brief.
- REQUIRED: enable the "include LinkedIn post text + hook" toggle.
- Export the top 50 (or however many the founder picked) as CSV to /campaigns/<date>/leads.csv.

2. PERSONALIZATION (Claude or ChatGPT, this prompt)

- For each row in the CSV, generate one 280-character DM.
- Reference the recommended hook (most recent post within 7 days).
- Pivot to the founder's company in one sentence with one concrete number from /facts/company-facts.md.
- Close with a low-friction ask.
- Save all 50 DMs to /campaigns/<date>/dms.md, one per partner,

numbered.

- PROMPT THE FOUNDER to read and approve.

3. DELIVERY (HeyReach MCP)

- Once the founder approves the batch, create a HeyReach campaign.
- Upload the 50 leads + 50 messages.
- Configure the 4-step sequence (day 0 / day 3 / day 7 / day 14).
- Set rate limits to LinkedIn safe defaults: 25 connections/day, 50 follow-ups/day, weekday only.
- Launch the campaign.

4. RESPONSE HANDLING (this prompt + HeyReach MCP)

- When HeyReach reports a reply, classify it (INTERESTED, INTERESTED-LATER, WRONG-FIT, WRONG-NOW, QUESTION).
- Draft a response under 320 characters in the same voice.
- Save the classification + draft to /campaigns/<date>/replies/<partner>.md.
- Update the partner's pipeline status via the Notion or Attio MCP.
- PROMPT THE FOUNDER to approve and send.

Refresh command

On day 7 of a running campaign, when the user types "refresh the campaign":

- Run an 8Raise partner-rotation check on all 50 LinkedIn URLs.
- For any partner who rotated firms, pause their sequence in HeyReach.
- Add the new partner contact at the new firm to a "next campaign" saved search.

Hard rules

- Never auto-send. Drafts only.
- Never exceed 25 connection requests per day.
- Never message the same partner twice in 6 months.
- Every DM under 280 characters.
- Every reply draft under 320 characters.
- Brand voice from /facts/founder-profile.md.

- All currency in EUR. All timestamps in CET.
- NEVER use em dashes. Use commas, colons, periods, or arrows.

Below the CLAUDE.md, the message templates.

Template: connection-request DM (280 chars max)

[Partner first name], your [post topic] thread mirrors what we saw building [Company]. [One-sentence pivot using the hook]. [One concrete number from /facts/]. [Low-friction ask: 15-min walkthrough / deck / single question].

Worked example (Ana García López, the Spanish fintech post):

Ana, your thread on SME-segment underserving in Spanish fintech mirrors we saw building Flowly. Mid-market SaaS teams sit in the same gap (47 integrations, no engineering team to wire them). €280k MRR from 18 customers since 2021. Would a 15-min walkthrough be useful?

Template: day-3 follow-up DM (after connection accepted)

Thanks for connecting, [Partner first name]. The reason I reached out: [one-line context]. We are seeing [one concrete metric or signal]. Quick deck attached / [calendar link]. Either works.

Template: day-7 different-angle DM

[Partner first name], one more data point that might be relevant to your [post topic / thesis]: [one specific customer or metric].

Worth a 15-min call?

Template: day-14 close-the-loop DM

[Partner first name], no worries if this is not a fit right now. Will close the loop on my end. Best of luck on [their fund's recent investment or initiative], that one looked interesting.

Template: response handler classification + draft

CLASSIFICATION: [INTERESTED / INTERESTED-LATER / WRONG-FIT / WRONG-NOW / QUESTION]

DRAFT (under 320 chars):

[Match their energy. Include the next concrete step (calendar link, deck share, single-paragraph answer to their question). No filler. Same brand voice as the original DM.]

Drop the **CLAUDE.md** plus the templates into a fresh Claude Code project. Run **claude** inside the folder. Type “run a 50-investor LinkedIn campaign for Flow European fintech raise” and you have a draft batch in 30 minutes and a live Heli campaign by lunch.

(INSERT IMAGE HERE: bonus illustration showing the CLAUDE.md file on the 5 message templates in the middle, and the four-layer execution flow on the post-08-claude-md.png)

Want us to run your LinkedIn outreach for you?

The four-layer loop above is the whole asset. It works if you wire it up yourself.

If you would rather not configure 8Raise, the AI prompts, HeyReach, the response handler, and the day-7 partner-rotation refresh, we run the full LinkedIn outreach as a service for a small number of founders preparing to raise.

→ We sit with you for the first build, run the 8Raise search against your real raise and tune the message prompt against your founder voice.

→ We set up the HeyReach campaign in your account, configure the 4-step sequence and connect rate limits to LinkedIn safe defaults.

→ We run the response handler with you for the first 14 days, classifying and categorizing replies until you have the rhythm.

→ We hand you the `CLAUDE.md`, the message templates, and the saved 8Raise so you can run the next campaign yourself.

→ You walk away with 50 active investor conversations, a typical 6 to 10 first calls booked, and the playbook to repeat the loop monthly.

We do not run 500 of these. We pick founders where we believe the loop will compound into a real round closing.

On the call we walk through your raise, look at your current outreach if you have it, and tell you honestly whether it is worth us running the loop with you or you running it yourself.

Hope this was valuable!

Cheers,

Nicolas Keller & Ignacio Prieto

8fundraising. From first outreach to funded.



9 Likes • 2 Restacks

Discussion about this post

Comments Restacks



Write a comment...